

Convolutional neuronal networks

Till Korten

With Material from

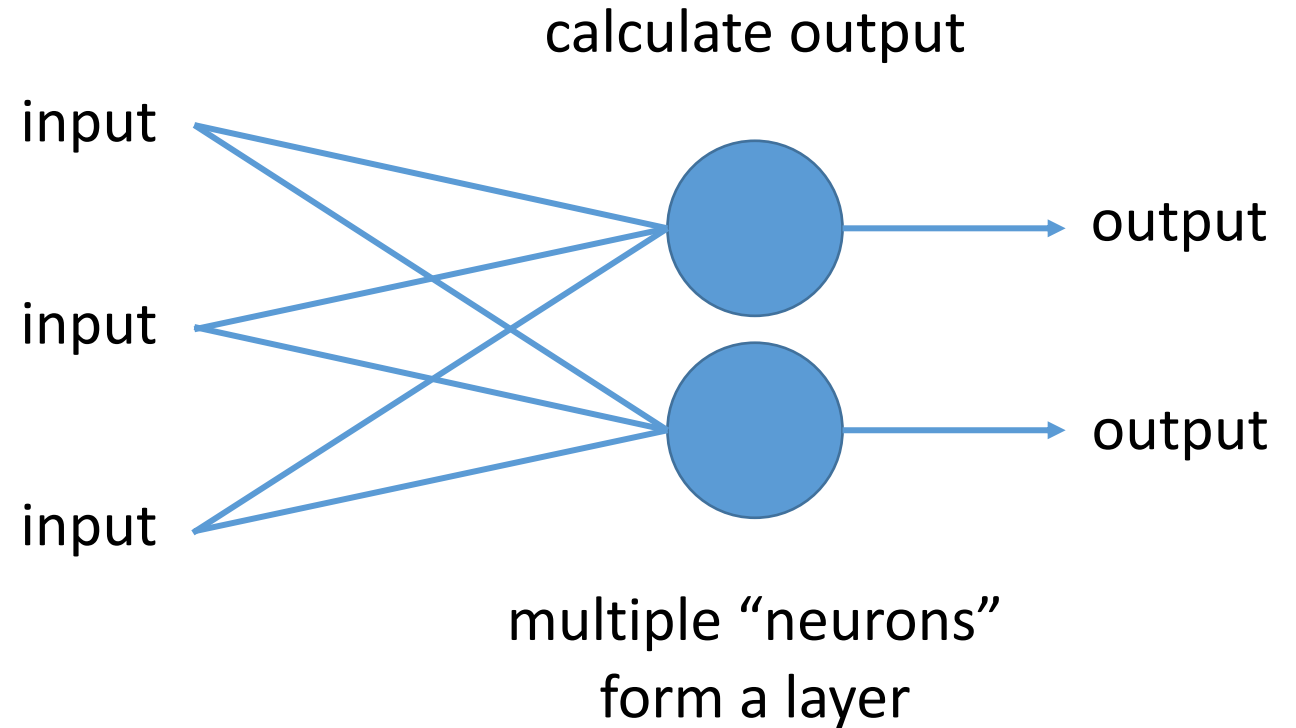

Johannes Müller, Robert Haase: PoL

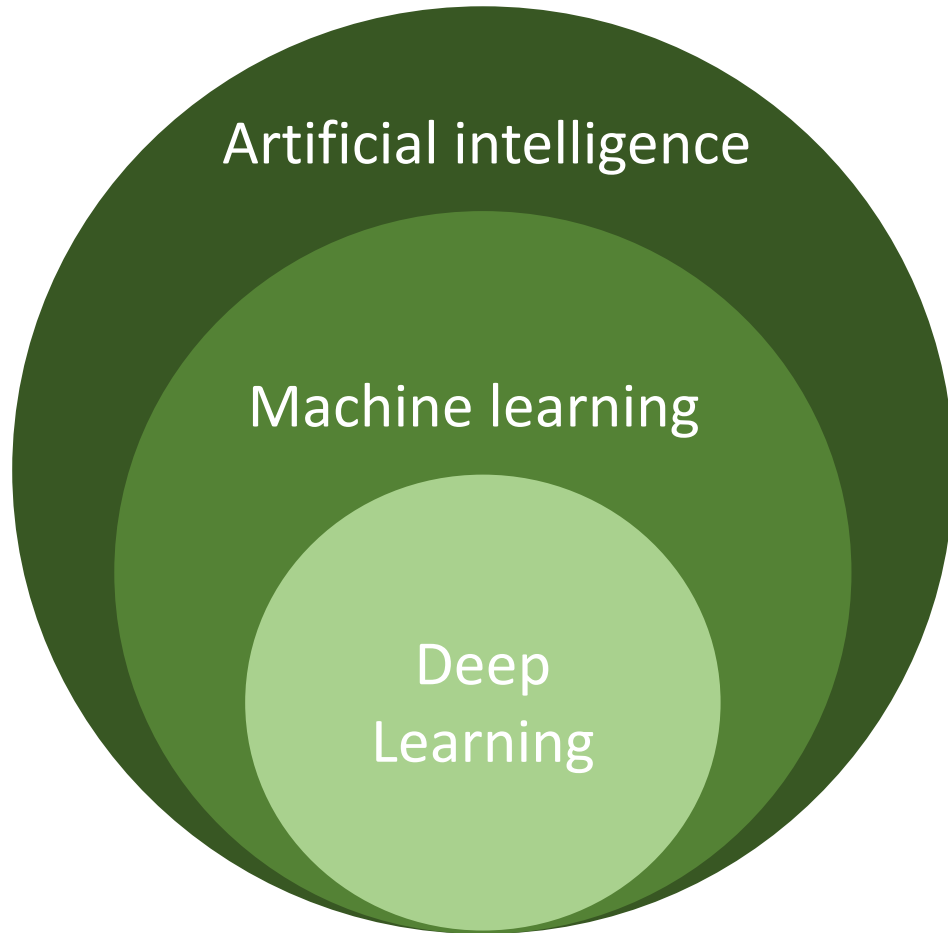
Alex Krull, Uwe Schmidt: MPI CBG

Martin Weigert: EPFL Lausanne


Ignacio Arganda-Carreras: Universidad del Pais Vasco

- Neural networks are composed of individual artificial “neurons”

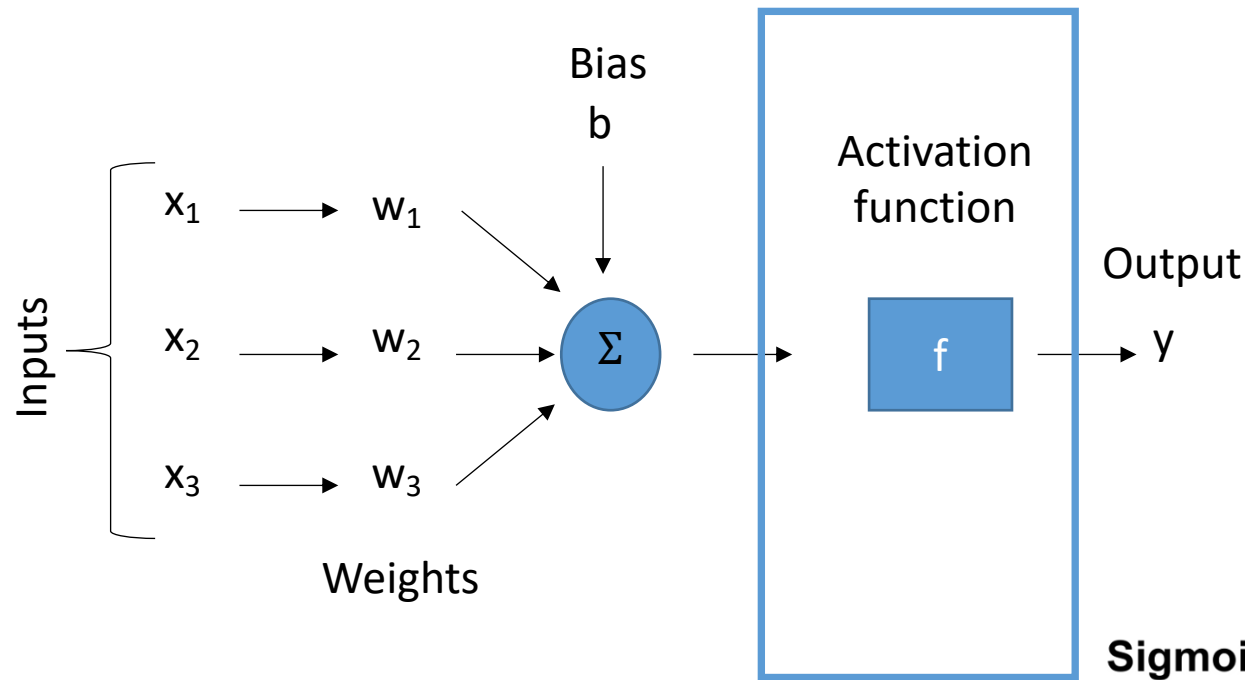




- Neural network: multiple layers of artificial “neurons”
- Deep neural network (DNN): neural network with more than one layer between input and output



An artificial “neuron” is a mathematical function



Single neuron output calculation

$$y = w_1 x_1 + w_2 x_2 + w_3 x_3 + b = w^T x + b$$

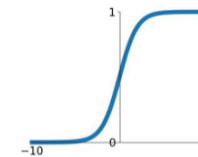
For image data, the values x_1, x_2, \dots would be

Pixel intensities

Pixel coordinates

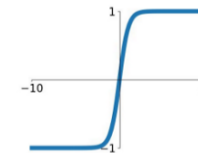
Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$



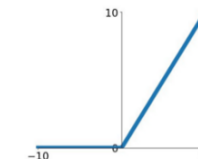
tanh

$$\tanh(x)$$




ReLU

$$\max(0, x)$$



Leaky ReLU

$$\max(0.1x, x)$$

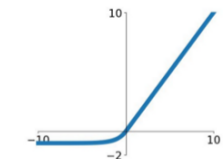


Maxout

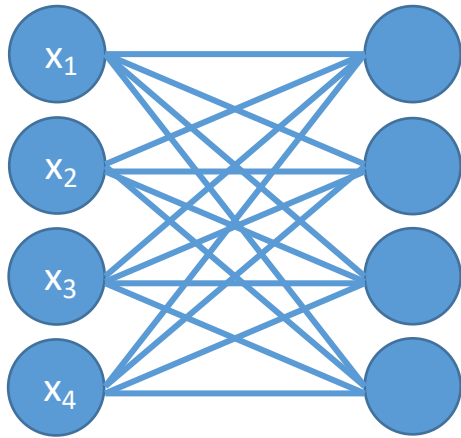
$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

ELU

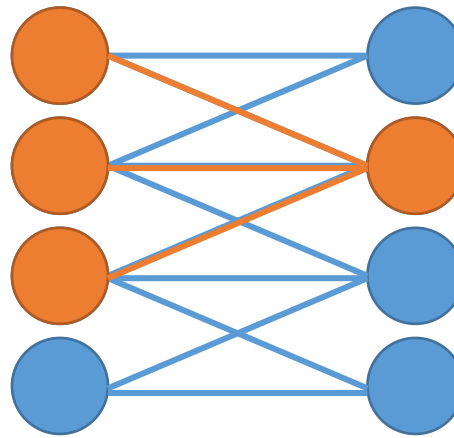
$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



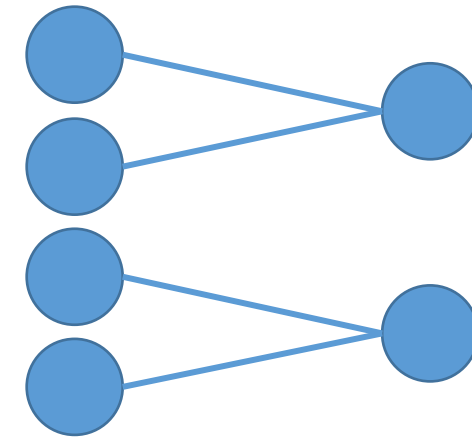
Network layers can have different architectures




Fully connected layer



Convolutional layer



Pooling layer

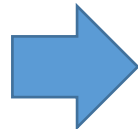


Convolutional layer

Previously:

Defined filter kernels

1/16	1/8	1/16
1/8	1/4	1/8
1/16	1/8	1/16

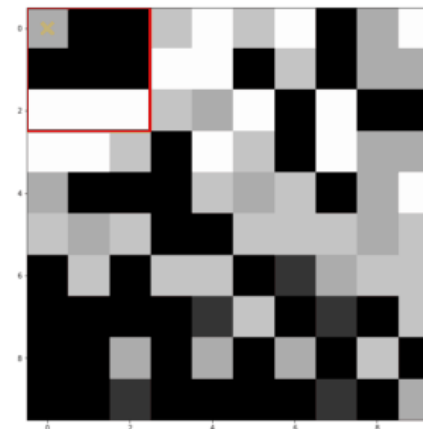


Now:

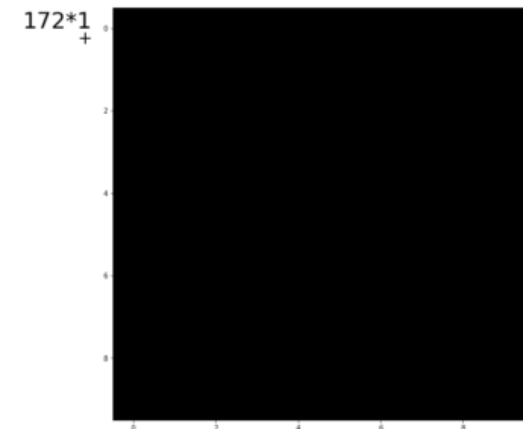
Learned filter kernels

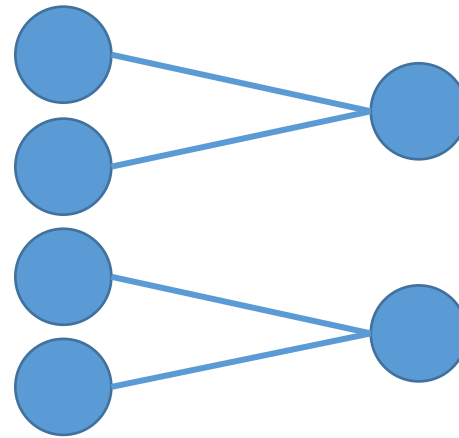
w_{11}	w_{12}	w_{13}
w_{21}	w_{22}	w_{23}
w_{31}	w_{32}	w_{33}

Input




Output





Pooling layer

3	15	1	13
9	7	0	10
11	5	5	3
1	8	9	6



Maximal values

15	13
11	9

Average values



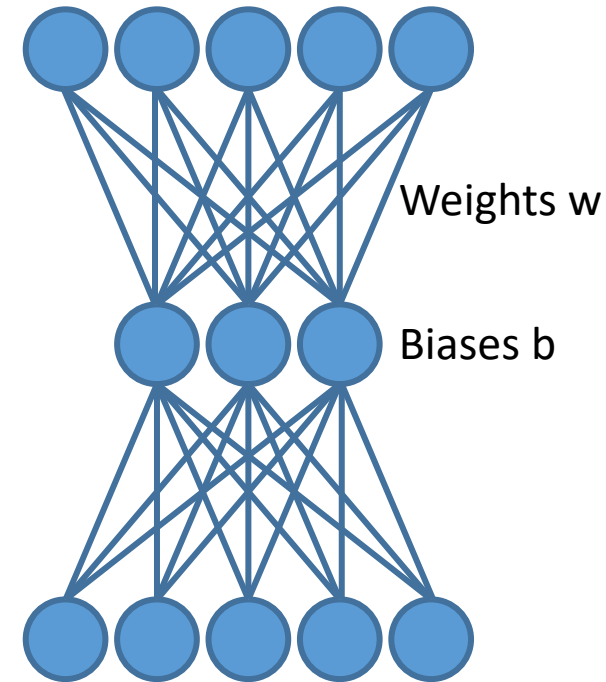
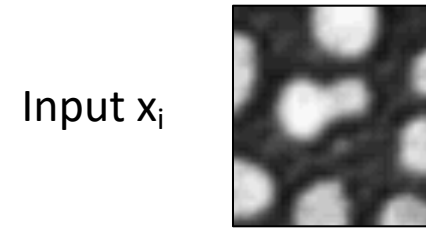
8.5	6.0
6.3	5.8

The network learns by minimizing a loss function

- Learning is an optimization problem
- Step 0: Initialize the network randomly
 - Weights
 - Bias
- Step 1: Forward pass the input through the network, get an initial prediction (Images 0...M)
- Step 2: Compare the output with the ground truth, compute the error (loss function)
 - The **loss function** can be freely defined.
 - Mean squared error:

$$\mathcal{L}(y, \hat{y}) = \frac{1}{M} \sum_{i=1}^M (\hat{y}_i - y_i)^2$$

- Step 3: Update weights



Ground truth y

The loss function can be expanded from

$$\mathcal{L}(y, \hat{y}) = \frac{1}{M} \sum_{i=1}^M (\hat{y}_i - y_i)^2$$




as the prediction depends on inputs x weights w and bias b

$$\mathcal{L}(\hat{y}, x, w) = \frac{1}{M} \sum_{i=1}^M ((w^T x_i + b) - y_i)^2$$

We can calculate derivatives with respect to w and b to find their optimal values

→ Derivatives tell us how to change w & b in order to improve the prediction (i.e. minimize the loss function)

Repeat this for each layer, update weights w



Popular frameworks:

<https://www.tensorflow.org/>

<https://www.pytorch.org/>

Hardware requirements: Nvidia (CUDA-capable) graphics card (GPU)

Memory: The more GPU memory the better



[Google LLC](#), Public domain, via Wikimedia Commons



[PyTorch](#), [BSD](#), via Wikimedia Commons



[Adam Kapetanakis](#), [CC BY-SA 4.0](#), via Wikimedia Commons

Before you start training:

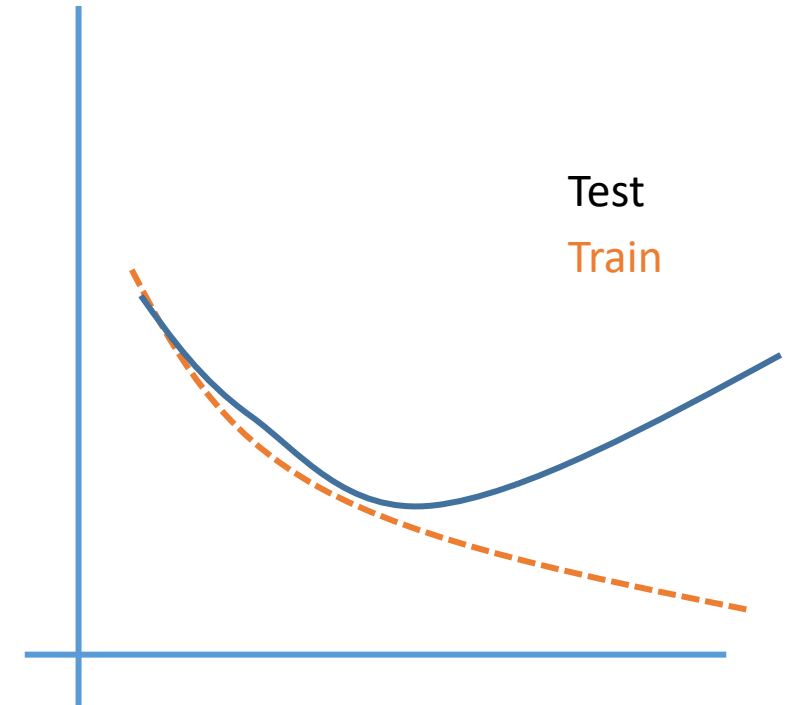
- Split data into three groups:
 - Training set
 - Testing set
 - Validation set

During training:

- Apply network to training set:
 - Measure performance (“loss”) of network, update weights
- Apply network to testing set:
 - Measure performance of updated network, keep weights unchanged
- Improve network architecture (optimize “hyperparameters”)

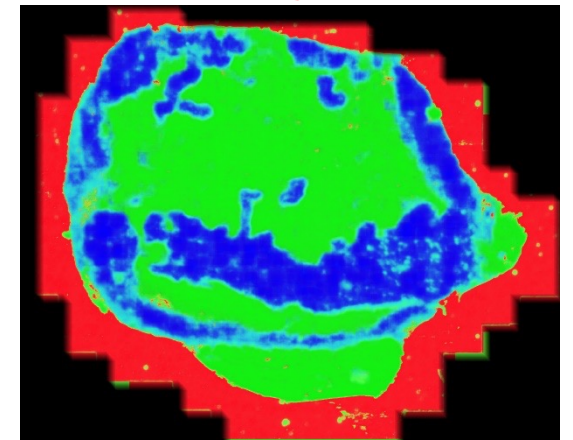
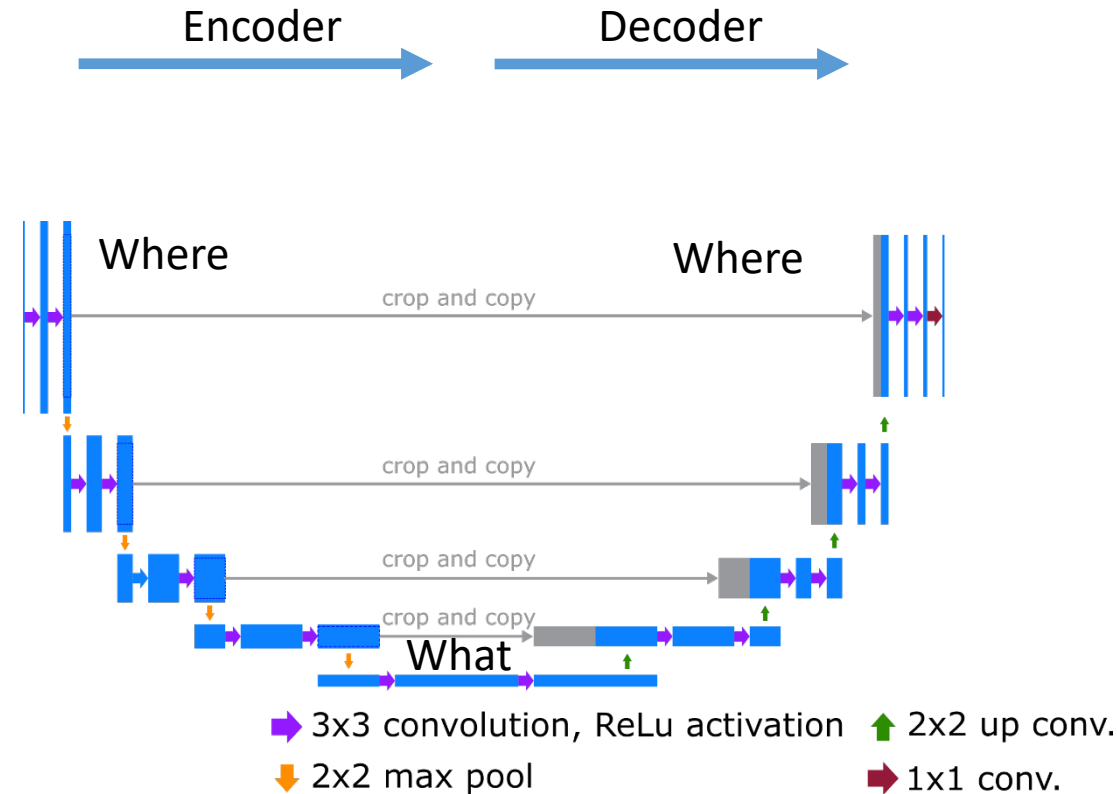
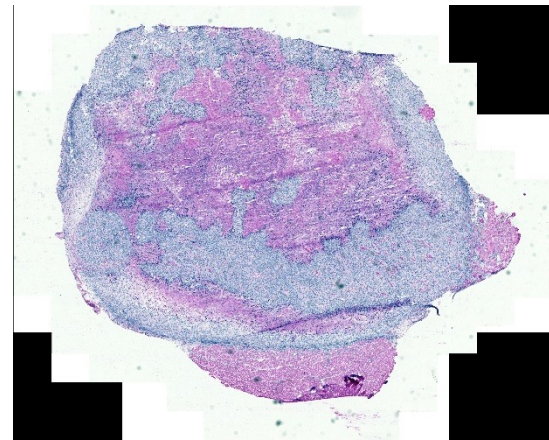
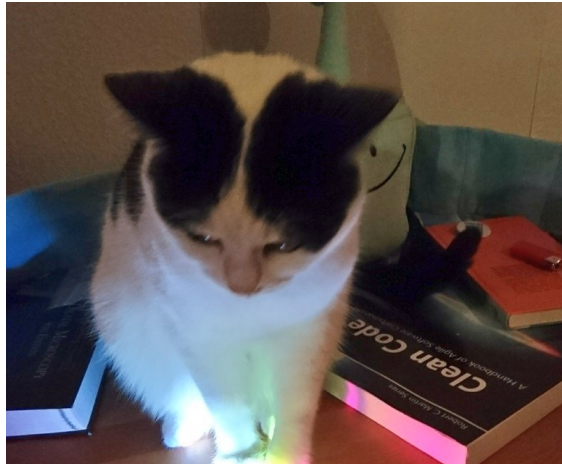
During validation:

- Measure performance in unseen validation dataset



Overfitting:





- Network is learning things „by heart“
- Hint at this happening: Updated weights from training fail to perform well in test



- The **U-net** is the most used network architecture in biological image processing using CNNs.
 - Encoder: Increase the “What”, decrease the “Where”
 - Decoder: Use the “What”, to identify the “Where”

Strategy:


- Add additional information to prediction
- Member pixels of objects (nuclei) can be reached via a straight line from the center




Dense Polygon Prediction
(e.g. U-Net, ResNet)

Polygon Selection
(Non-Maximum Suppression NMS)

Noisy images + Crowded cells = Common source of segmentation errors




Dense Segmentation
(e.g. U-Net)



Bounding box based methods
(e.g. Mask-RCNN)

Object probabilities

Radial Distances



© Schmid et al., MICCAI (2018)

Problem:

- Multiple candidate points for nucleus center
- Overlapping instance predictions

Before NMS



After NMS





Non-maximum-suppression (NMS):

- Intersection over Union (IoU) threshold τ determines „conservativeness“:
 - High τ : Objects tend to be considered as separate objects
 - Low τ : Objects tend to be considered as the same objects

Object probabilities

Radial Distances











© Schmid et al., MICCAI (2018)


Non-maximum-suppression (NMS):

- Object probabilities: Probability that pixel belongs to class “nucleus”
- Multiple maxima lead to multiple possible polygons for the same nucleus

Algorithm:

- Select polygon with highest object probability inside: 
- Look at other polygons: Is the overlap of  with  larger than threshold τ ?
 - Yes:  and  are actually the same object, drop 
 - No:  and  are separate nuclei
- Setting τ very high leads to many false positives!



Non-maximum suppression



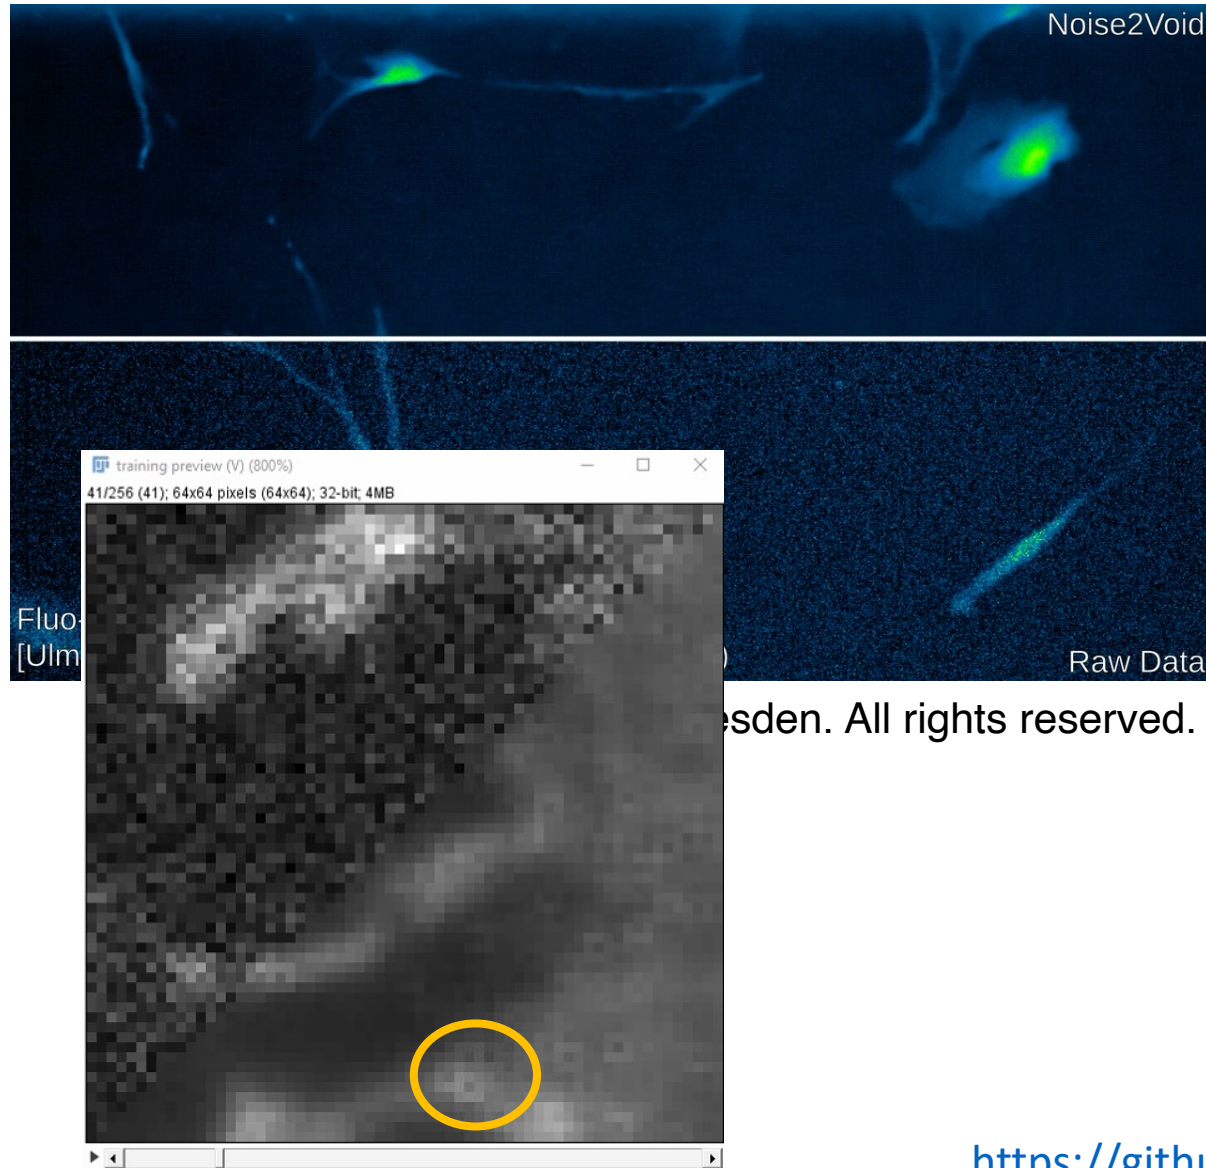
© Schmid et al., MICCAI (2018)

- **CARE**: content-aware restoration
- Image acquisition of pairs of images: A high-quality and a low-quality image.
- Caveats:
 - Reconstructs shot noise present in high quality training images
 - Trained model only applicable to image data of the same conditions (biological sample, microscope, etc)


5 example validation patches
top row: input (source), middle row: target (ground truth), bottom row: predicted from source



© MPI-CBG, Dresden. All rights reserved.



Raw data = signal + noise



Strategy:

- Try to predict intensity of pixel y from surrounding pixels x
- CNN fails to predict noise component → N2V can only reproduce signal from the surroundings of y

Beware:

- Only **random** noise can be removed, otherwise artifacts occur

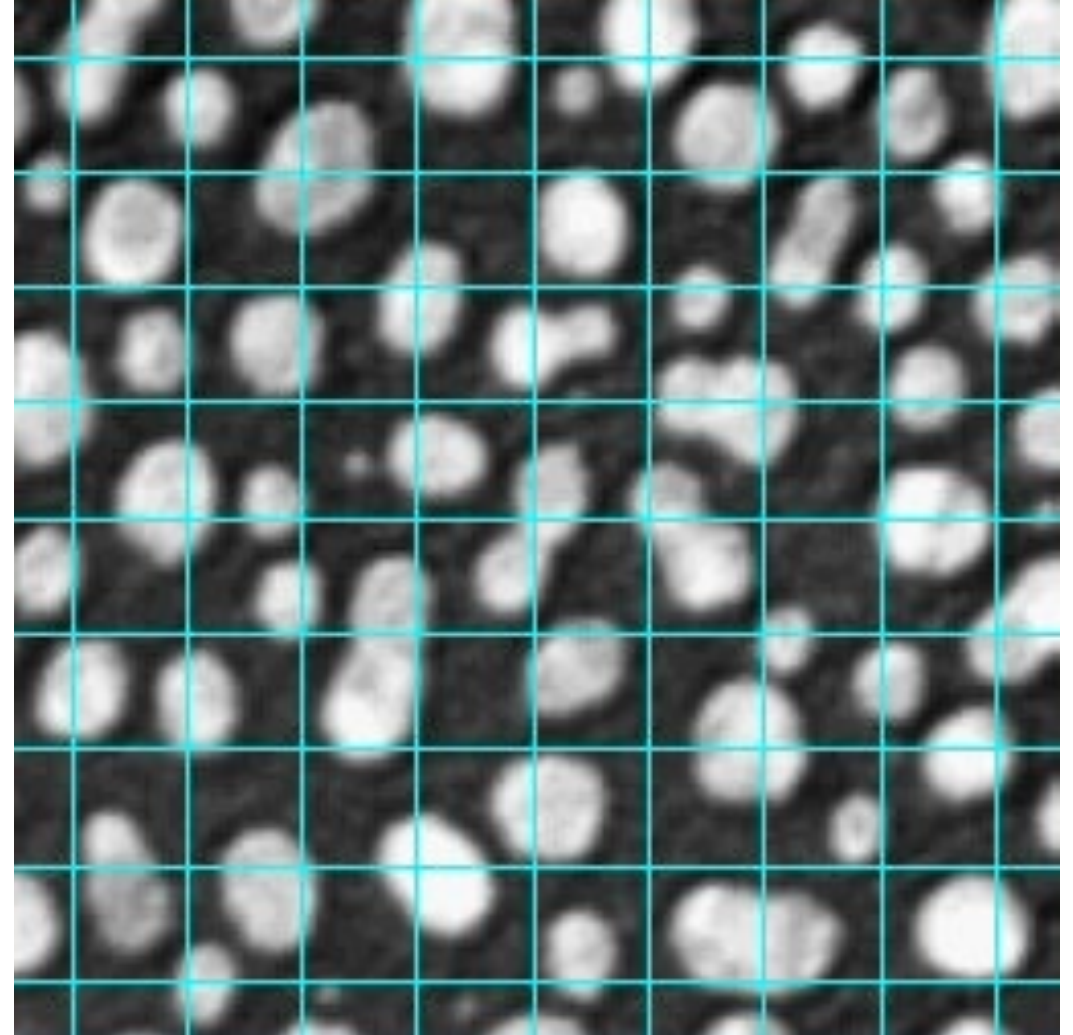
<https://github.com/juglab/n2v>

<https://forum.image.sc/t/n2v-artefacts-in-training-data/70686>

- Images are tiled
- limited “receptive field of the network”

Receptive field:

- Objects must be smaller than receptive field to be detectable



Unbalanced training data:

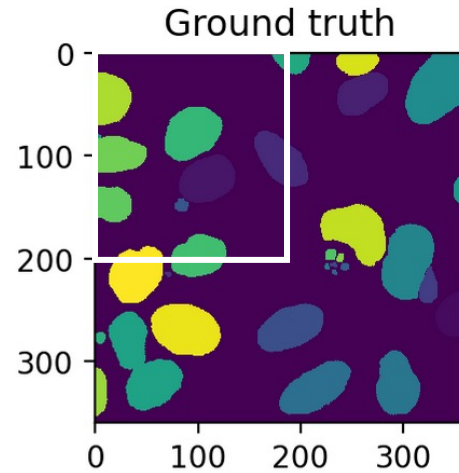
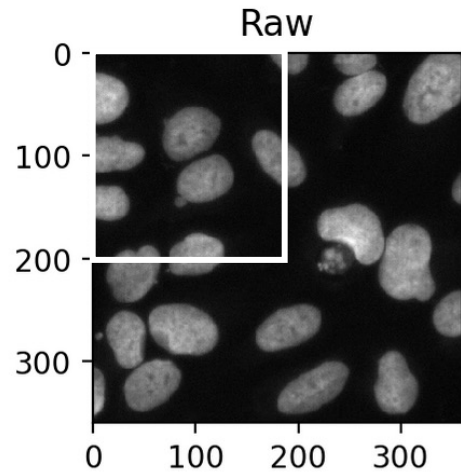
- Some labels appear more often in training data than others
- Rare events will not be learned because missing them doesn't harm accuracy much
- Weighted data sampling
- **Biased results!!**

Is the iPhone racist? Chinese users claim iPhoneX face recognition can't tell them apart


APPLE has come under fire following numerous complaints from Chinese users who claim the iPhone X face recognition can't tell them apart.

<https://www.news.com.au/technology/gadgets/mobile-phones/is-the-iphone-racist-chinese-users-claim-iphonex-face-recognition-cant-tell-them-apart/news-story/13814540e8c82ad466aca687e12af64c>

When the input data does not fit to the training data



Prediction


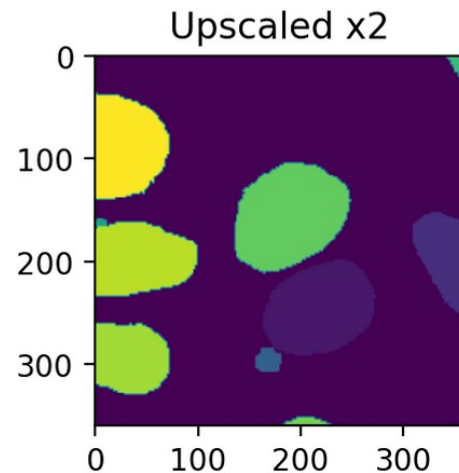
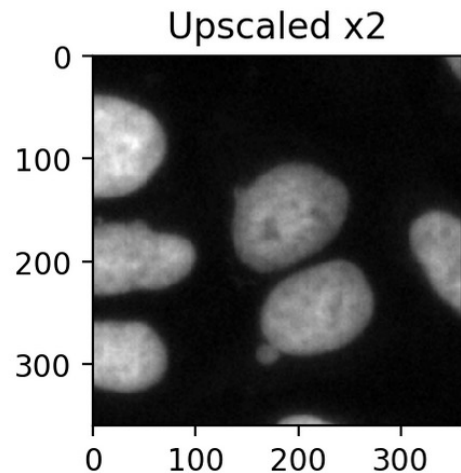


What happened here?

Receptive field too small

I used a different resolution than during training

Overfitting



- With great power comes great responsibility: **Validate your models well!**
- Better data more important than better model
- Often performs fantastic – *but you don't know why*
- Generative neural networks (like CARE) can dream up data – to a hammer everything looks like a nail!